

# Arabic Sign Language Real-Time Recognition

(Using YOLOv8 Object Detection)



# Table of contents

**01**

Arabic Sign  
Language (ASL)

**02**

YOLO Architecture

**03**

Dataset

**04**

Model Fine-Tuning

**05**

Metrics & Results

**06**

Code Snippets



**01**

# **Arabic Sign Language (ASL)**

# Arabic Sign Language



































## Arabic Sign Language (ASL)

A visual-gestural language used by deaf and hard of hearing communities primarily in our Arabian world. Just like spoken languages, ASL has its own grammar, syntax, and vocabulary. It involves the use of handshapes, movements to convey meaning and communicate ideas. In recent years, there has been increased attention to developing technologies that facilitate the recognition and translation of ASL gestures, aiming to enhance accessibility and communication for individuals who use sign language.



Lam - ل

# Arabic Sign Language

ألف - Alif	باء - Bā	آاء - Tā	ثاء - Thā	جيم - Jīm	حاء - Hā	خاء - Khā	دال - Dāl
							
ذال - Dhāl	راء - Rā	زاي - Zāy	سين - Sīn	شين - Shīn	صاد - Sād	داد - Dād	طاء - Tā
							
ظاء - Zā	عين - Ayn	غين - Ghayn	فاء - Fā	قاف - Qāf	كاف - Kāf	لام - Lām	ميم - Mīm
							
نون - Nūn	هاء - Hā	واو - Wāw	ياء - Yā	ة - Tāa	ال - Al	لا - Laa	ياء - Yāa
							

Arabic Letters resembled by sign language

# YOLO Architecture



## YOLO (You Only Look Once)

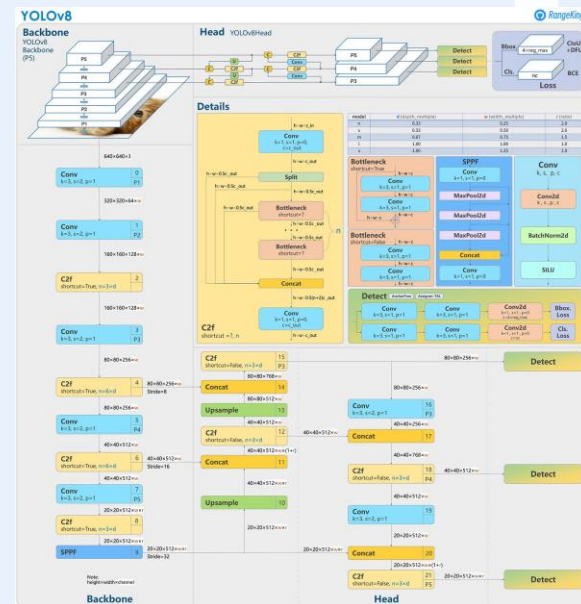
YOLO is known as a SOTA (State-Of-The-Art) approach and architecture that's widely used in object detection along with many other tasks. We used the YOLOv8 model

YOLO consists of many parts, mainly:

**1. Backbone:** The backbone is the foundational part of the architecture responsible for extracting meaningful features from the input image, typically made of Convolutional Neural Networks (CNNs) such as DarkNet in YOLO.

**2. Neck:** The neck sits between the backbone and the head. It performs additional feature fusion or transformation to enhance the representations learned by the backbone.

**3. Head:** the head consists of convolutional layers that predict bounding boxes, objectness scores (indicating the presence of an object), and class probabilities for each grid cell in the feature map generated by the backbone and neck.



YOLOv8Architecture


# Dataset



## Arabic Sign Language Dataset 2022

<https://www.kaggle.com/datasets/ammarsayedtaha/arabic-sign-language-dataset-2022>

The Dataset consists of 14202 images covering the 32 Arabic Signs of letters by hands made by 50 different persons, 9955 samples are used for training, while the rest 4247 samples are used for evaluating the model for generalization and performance on unseen data.



9 Plotted samples with their corresponding boundary box and label

# Model Fine-Tuning



## Hyper-Parameters:

*Epochs: 100*

*Image Size: 256*

*batch Size: 128*


*Optimizer: Adam*

*Initial Learning Rate: 0.001*

*momentum (beta for Adam): 0.9*

*Learning Rate Scheduler: Cosine Scheduler*

*Random Seed: 42*



**Cosine Learning Rate Scheduler  
Effect**



# Model Fine-Tuning



Model's Parameters:



3.017 Million Parameters



8.227 GFLOPs

# Metrics & Results



## System Specs:

*GPU: Nvidia RTX 3060 12GB VRAM*

*CPU: Intel I5-10400F 2.90 GHz*

*RAM : 16GB*

*OS: Windows 11 Pro*



## Frameworks:

- *Visual Studio Code*
- *Python*
- *Ultralytics (YOLO)*
- *OpenCV*




# Metrics & Results



## mAP50

measures the average precision of correctly identifying background regions or areas with no objects of interest in the image. It evaluates how well the model can distinguish between objects and background regions.




**mAP50 Through the epochs**  
**98.24% at Epoch 100**

# Metrics & Results



## mAP50-95

calculates the average precision across a range of IoU thresholds (from 0.5 to 0.95) for background regions. It provides a broader assessment of the model's performance across various levels of overlap between predicted and ground truth bounding boxes.




**mAP50-95 Through the epochs**  
**81.89% at Epoch 100**

# Metrics & Results



## Precision

measures the proportion of correctly predicted background regions out of all regions classified as background by the model. It indicates the model's ability to avoid falsely identifying objects in background regions.




**Precision Through the epochs**  
**96.14% at Epoch 100**

# Metrics & Results





## Recall

calculates the proportion of correctly identified background regions out of all actual background regions in the image. It assesses the model's ability to detect all background regions, including those that might be challenging to differentiate from objects.

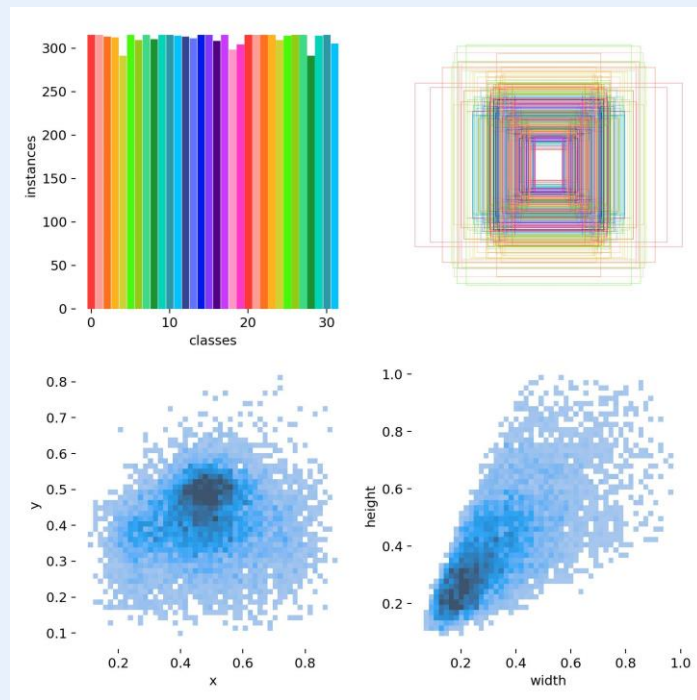


Precision Through the epochs  
96.14% at Epoch 100

# Metrics & Results





# Metrics & Results






# Metrics & Results

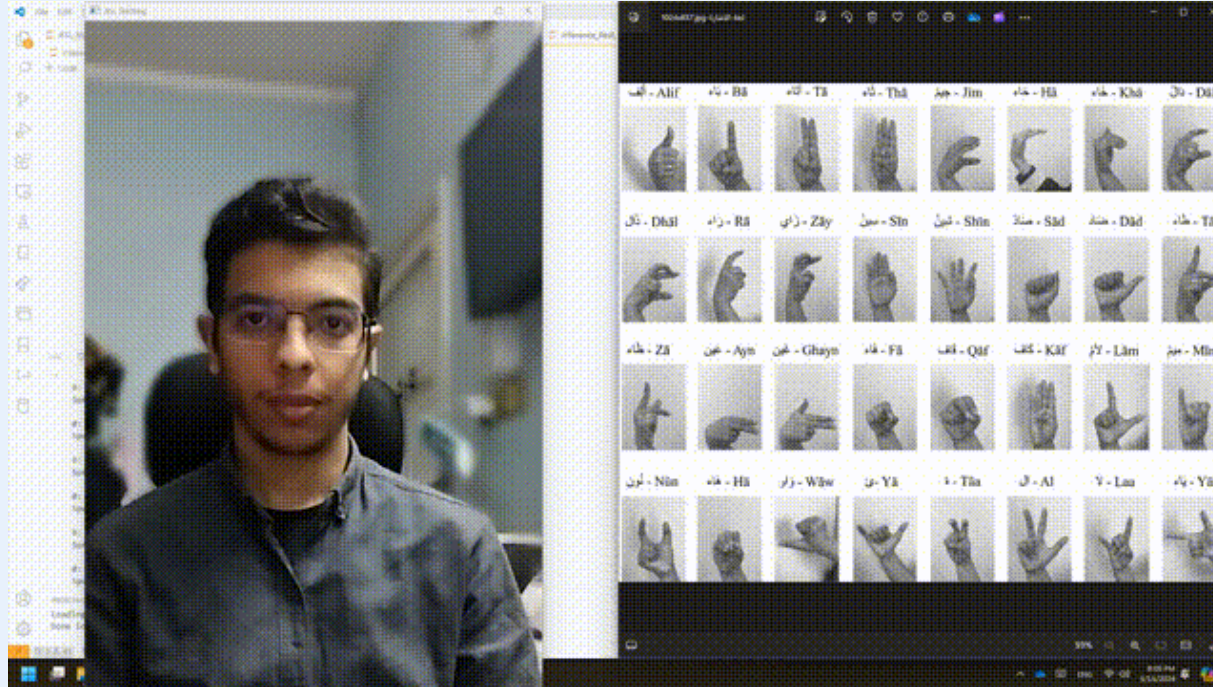


# Metrics & Results



Inference of the model on some images of the Validation set

# Metrics & Results



A Real-Time Inference of the model

# Code Snippets



```
1 from ultralytics import YOLO
2 import os
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import cv2
6 import random
7
8 import comet_ml
9 comet_ml.init()
```

Imports

# Code Snippets

```
1 def read_labels(label_path):
2     with open(label_path, 'r') as file:
3         lines = file.readlines()
4         labels = [line.strip().split() for line in lines]
5         return labels
6
7 def draw_boxes(image, labels):
8     for label in labels:
9         class_id = int(label[0])
10        x, y, w, h = map(float, label[1:])
11        image_height, image_width, _ = image.shape
12        x1 = int((x - w / 2) * image_width)
13        y1 = int((y - h / 2) * image_height)
14        x2 = int((x + w / 2) * image_width)
15        y2 = int((y + h / 2) * image_height)
16        color = (0, 255, 0) # Green
17        cv2.rectangle(image, (x1, y1), (x2, y2), color, 2)
18    return image
19
20 data_dir = 'train'
21
22 image_files = [file for file in os.listdir(os.path.join(data_dir, 'images')) if file.endswith('.jpg')]
23
24 random.shuffle(image_files)
25
26 fig, axes = plt.subplots(3, 3, figsize=(12, 12))
27 for ax, image_file in zip(axes.ravel(), image_files[:9]):
28     image_path = os.path.join(data_dir, 'images', image_file)
29     label_path = os.path.join(data_dir, 'labels', os.path.splitext(image_file)[0] + '.txt')
30
31     # Read image
32     image = cv2.imread(image_path)
33     image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
34
35     # Read Labels
36     labels = read_labels(label_path)
37
38     # Draw bounding boxes on the image
39     image_with_boxes = draw_boxes(image_rgb, labels)
40
41     # Display image with bounding boxes and set the labels as titles
42     ax.imshow(image_with_boxes)
43     ax.set_title(str(labels[0][0]))
44     ax.axis('off')
45
46 plt.tight_layout()
47 plt.show()
```

Visualizing Training Samples

# Code Snippets



```
1 model = YOLO('yolov8n.yaml')  
2 model = YOLO('yolov8n.pt')  
3 model = YOLO('yolov8n.yaml').load('yolov8n.pt')
```



```
1 history = model.train(data='sign.yaml', epochs=100, imgsz=256,  
2                       patience = 100, batch = 128,  
3                       project = "ASL", optimizer = 'Adam', momentum = 0.9,  
4                       cos_lr=True, seed = 42, plots = True, close_mosaic = 0, lr0 = 0.001)
```

## Model Training

# Code Snippets



```
1 # Loading Model
2 trained_model = YOLO('ASL.pt')
```



```
1 test_images_dir = 'valid/images'
2 test_images = [os.path.join(test_images_dir, image) for image in os.listdir(test_images_dir)]
3
4 test_samples = np.random.choice(test_images, 10, replace=False)
5
6 print(test_samples)
```

## Loading Model and Making Predictions

# Code Snippets

```
1 results = trained_model([test_samples[0], test_samples[1], test_samples[2], test_samples[3], test_samples[4], test_samples[5], test_samples[6], test_samples[7], test_samples[8], test_samples[9]])
2
3 results_dir = "results_tries"
4 if not os.path.exists(results_dir):
5     os.makedirs(results_dir)
6
7 i = 0
8 for result in results:
9     boxes = result.boxes
10    masks = result.masks
11    keypoints = result.keypoints
12    probs = result.probs
13
14    # Save image
15    filename = os.path.join(results_dir, f"result_{i}.jpg")
16    result.save(filename=filename)
17    # result.show() # Uncomment to display image
18    i+=1
19
```

## Loading Model and Making Predictions



# Code Snippets



```
1 directory = "results_tries"
2
3 images = []
4 for filename in os.listdir(directory):
5     if filename.endswith(".jpg"):
6         img = plt.imread(os.path.join(directory, filename))
7         images.append(img)
8
9 fig, axs = plt.subplots(2, 5, figsize=(50, 25))
10 for i in range(10):
11     axs[i//5, i%5].imshow(images[i])
12     axs[i//5, i%5].axis('off')
13
14 plt.show()
```

Loading Model and Making Predictions

# Code Snippets

```
1 import cv2
2 from ultralytics import YOLO
3
4 # Load the model
5 model = YOLO('ASL.pt')
6
7 # Open the camera
8 cap = cv2.VideoCapture(0)
9
10 # Setting width and height of the video
11 cap.set(cv2.CAP_PROP_FRAME_WIDTH, 1920)
12 cap.set(cv2.CAP_PROP_FRAME_HEIGHT, 1080)
13
14
15 while cap.isOpened():
16
17     success, frame = cap.read()
18
19     if success:
20
21         results = model.track(frame, persist=True)
22
23         annotated_frame = results[0].plot()
24
25         cv2.imshow("ASL Tracking", annotated_frame)
26
27         # Break the Loop if 'q' is pressed
28         if cv2.waitKey(1) & 0xFF == ord("q"):
29             break
30     else:
31         # Break the Loop if
32         break
33
34 cap.release()
35 cv2.destroyAllWindows()
36
37
```

## Real-Time Inference

**Thank You!**

